

A1 Report

February 4, 2025 4:05 PM

Benny Liang | 3019242

Q1)

a) Use the time utility to time palindrome.py and slow-pali.cpp on files t3.txt and t4.txt. Copy/paste the output of time from the terminal window into your report.

Using t3.txt as input

	`palindrome.py`	`slow-pali.cpp`
Real	0m0.034s	0m0.012s
User	0m0.022s	0m0.004s
Sys	0m0.011s	0m0.006s

Using t4.txt as input

	`palindrome.py`	`slow-pali.cpp`
Real	0m0.249s	0m2.295s
User	0m0.233s	0m0.361s
Sys	0m0.013s	0m1.908s

b) Using your results from (a), how much time did the C++ and python programs spend executing on the CPU, and how much time did each of them spend waiting for I/O to finish?

CPU time spent executing on `t3.txt`

- `palindrome.py` spent $0.022s + 0.011s = 0m0.033s$ on the CPU
- `slow-pali.cpp` spent $0.004s + 0.006s = 0m0.010s$ on the CPU

CPU time spent executing on `t4.txt`

- `palindrome.py` spent $0.233s + 0.013s = 0m0.246s$ on the CPU
- `slow-pali.cpp` spent $0.361s + 1.908s = 0m2.269s$ on the CPU

Time spent waiting for I/O to finish on `t3.txt`

- `palindrome.py` I/O time taken = $real - user - sys = 0.034s - 0.022s - 0.011s = 0m0.001s$
- `slow-pali.cpp` I/O time taken = $real - user - sys = 0.012s - 0.004s - 0.006s = 0m0.002s$

Time spent waiting for I/O to finish on `t4.txt`

- `palindrome.py` I/O time taken = $real - user - sys = 0.249s - 0.233s - 0.013s = 0m0.003s$
- `slow-pali.cpp` I/O time taken = $real - user - sys = 2.295s - 0.361s - 1.908s = 0m0.026s$

c) Run 'strace -c' on palindrome.py and slow-pali.cpp on t3.txt and t4.txt. Copy/paste the output from the terminal window into your report.

`palindrome.py`

`t3.txt` as input						`t4.txt` as input					
Longest palindrome: ____o.O.o____						Longest palindrome: redder					
% time	seconds	usecs/call	calls	errors	syscall	% time	seconds	usecs/call	calls	errors	syscall
16.94	0.000728	10	71	20	openat	20.55	0.000995	199	5	3	execve
12.73	0.000547	109	5	3	execve	15.20	0.000736	15	49		mmap
12.08	0.000519	10	49		mmap	14.09	0.000682	9	71	20	openat
9.07	0.000390	2	141	55	newfstatat	10.45	0.000506	3	141	55	newfstatat
8.63	0.000371	5	73		fstat	9.33	0.000452	6	70		read
7.19	0.000309	6	51		close	6.98	0.000338	4	73		fstat
6.82	0.000293	6	48		read	6.05	0.000293	5	51		close

6.58	0.000283	15	18	getdents64	2.91	0.000141	7	18	getdents64
6.42	0.000276	4	66	rt_sigaction	2.79	0.000135	3	45	2 lseek
3.05	0.000131	2	45	2 lseek	2.38	0.000115	14	8	mprotect
2.77	0.000119	10	11	brk	1.45	0.000070	2	24	19 ioctl
1.98	0.000085	3	24	19 ioctl	1.28	0.000062	5	11	brk
1.49	0.000064	8	8	mprotect	1.22	0.000059	11	5	munmap
0.86	0.000037	7	5	munmap	0.89	0.000043	10	4	pread64
0.63	0.000027	9	3	getrandom	0.74	0.000036	0	66	rt_sigaction
0.35	0.000015	3	4	fcntl	0.74	0.000036	18	2	2 access
0.33	0.000014	7	2	futex	0.41	0.000020	10	2	arch_prctl
0.26	0.000011	5	2	rt_sigprocmask	0.41	0.000020	10	2	set_robust_list
0.26	0.000011	2	4	pread64	0.41	0.000020	10	2	rseq
0.23	0.000010	5	2	getcwd	0.39	0.000019	9	2	set_tid_address
0.19	0.000008	4	2	2 access	0.35	0.000017	5	3	getrandom
0.14	0.000006	1	4	3 readlink	0.25	0.000012	3	4	3 readlink
0.14	0.000006	6	1	gettid	0.25	0.000012	6	2	futex
0.12	0.000005	2	2	arch_prctl	0.25	0.000012	6	2	prlimit64
0.12	0.000005	2	2	set_tid_address	0.12	0.000006	1	4	fcntl
0.12	0.000005	2	2	set_robust_list	0.10	0.000005	2	2	getcwd
0.12	0.000005	2	2	prlimit64	0.00	0.000000	0	1	write
0.12	0.000005	2	2	rseq	0.00	0.000000	0	2	rt_sigprocmask
0.09	0.000004	4	1	write	0.00	0.000000	0	1	getuid
0.07	0.000003	3	1	getgid	0.00	0.000000	0	1	getgid
0.05	0.000002	2	1	getuid	0.00	0.000000	0	1	geteuid
0.05	0.000002	2	1	geteuid	0.00	0.000000	0	1	getegid
0.05	0.000002	2	1	getegid	0.00	0.000000	0	1	gettid
-----					-----				
100.00	0.004298	6	654	104 total	100.00	0.004842	7	676	104 total

`slow-pali.cpp`

`t3.txt` as input					`t4.txt` as input				
Longest palindrome: __o.O.o__					Longest palindrome: redder				
% time	seconds	usecs/call	calls	errors syscall	% time	seconds	usecs/call	calls	errors syscall

37.90	0.000669	669	1	execve	99.99	18.229164	3	5767198	read
23.34	0.000412	17	23	mmap	0.00	0.000580	580	1	execve
14.96	0.000264	6	43	read	0.00	0.000560	24	23	mmap
7.20	0.000127	21	6	mprotect	0.00	0.000099	19	5	openat
4.31	0.000076	15	5	openat	0.00	0.000068	11	6	mprotect
2.61	0.000046	7	6	fstat	0.00	0.000057	9	6	fstat
2.04	0.000036	7	5	close	0.00	0.000053	10	5	close
1.76	0.000031	31	1	munmap	0.00	0.000037	12	3	brk
1.42	0.000025	12	2	pread64	0.00	0.000028	14	2	pread64
0.74	0.000013	4	3	brk	0.00	0.000019	19	1	1 access
0.62	0.000011	11	1	arch_prctl	0.00	0.000012	12	1	set_tid_address
0.62	0.000011	11	1	set_tid_address	0.00	0.000012	12	1	rseq
0.62	0.000011	11	1	set_robust_list	0.00	0.000011	11	1	write
0.62	0.000011	11	1	rseq	0.00	0.000011	11	1	arch_prctl
0.57	0.000010	10	1	write	0.00	0.000011	11	1	set_robust_list
0.45	0.000008	8	1	prlimit64	0.00	0.000009	9	1	getrandom
0.23	0.000004	4	1	futex	0.00	0.000000	0	1	munmap
0.00	0.000000	0	1	1 access	0.00	0.000000	0	1	futex
0.00	0.000000	0	1	getrandom	0.00	0.000000	0	1	prlimit64
-----					-----				
100.00	0.001765	16	104	1 total	100.00	18.230731	3	5767259	1 total

d) Using the results from (c), why is the python program faster on some inputs than the C++ program, but slower on others?

The python program was slower on *smaller inputs* due to the fact that it had more read calls than the C++ program on the same input. Another contributor are the types of languages Python and C++ are, Python is an

interpreted language (runs line-by-line) whereas C++ is compiled before execution, helping the C++ program run faster than the Python program.

The python program was faster than the C++ program on *very large inputs*, due to the fact that the C++ program had to make so many `read` calls. When comparing with the python program it had a fraction of the C++ program's read calls.

Q3)

a) Run your fast-pali.cpp on t3.txt and t4.txt files using 'time' and 'strace -c'. Copy/paste the output from the terminal window into your report.

`t3.txt` as input

`time`	`strace -c`
Longest palindrome: ____o.O.o____	Longest palindrome: ____o.O.o____
real 0m0.009s	% time seconds usecs/call calls errors syscall
user 0m0.004s	-----
sys 0m0.004s	52.14 0.000573 24 23 mmap
	18.65 0.000205 34 6 mprotect
	9.01 0.000099 19 5 openat
	4.64 0.000051 8 6 fstat
	4.55 0.000050 8 6 read
	4.09 0.000045 9 5 close
	2.55 0.000028 14 2 pread64
	1.18 0.000013 13 1 arch_prctl
	1.09 0.000012 12 1 set_tid_address
	1.09 0.000012 12 1 set_robust_list
	1.00 0.000011 11 1 rseq
	0.00 0.000000 0 1 write
	0.00 0.000000 0 1 munmap
	0.00 0.000000 0 3 brk
	0.00 0.000000 0 1 1 access
	0.00 0.000000 0 1 execve
	0.00 0.000000 0 1 futex
	0.00 0.000000 0 1 prlimit64
	0.00 0.000000 0 1 getrandom

	100.00 0.001099 16 67 1 total

`t4.txt` as input

`time`	`strace -c`
Longest palindrome: redder	Longest palindrome: redder
real 0m0.111s	% time seconds usecs/call calls errors syscall
user 0m0.071s	-----
sys 0m0.037s	31.24 0.001987 180 11 read
	30.58 0.001945 97 20 brk
	11.26 0.000716 23 31 mmap
	10.34 0.000658 658 1 execve
	9.50 0.000604 67 9 munmap
	2.15 0.000137 27 5 openat
	1.48 0.000094 18 5 close
	1.15 0.000073 12 6 fstat
	0.61 0.000039 19 2 pread64
	0.33 0.000021 21 1 1 access
	0.30 0.000019 19 1 set_tid_address
	0.25 0.000016 16 1 set_robust_list
	0.20 0.000013 13 1 futex
	0.20 0.000013 13 1 getrandom
	0.19 0.000012 12 1 arch_prctl
	0.19 0.000012 12 1 rseq
	0.03 0.000002 2 1 write

0.00	0.000000	0	6	mprotect
0.00	0.000000	0	1	prlimit64

100.00	0.006361	60	105	1 total

b) Is your fast-pali.cpp faster than slow-pali.cpp? Why do you think that is?

Yes my `fast-pali.cpp` is much faster than `slow-pali.cpp` because of the significantly reduced number of system calls for `read`. For the `t4.txt` input; in `slow-pali.cpp` it had **5756198** read calls while `fast-pali.cpp` had **11** read calls, significant difference. Then looking at the `time` results (for `t4.txt`), `fast-pali.cpp` finished much faster than `slow-pali.cpp`, again due to the significantly reduced number of `read` calls.

C) Is your program faster than palindrome.py and why?

Yes `fast-pali.cpp` is faster than `palindrome.py` for similar reasons as described in b). The results for `t4.txt` input, `palindrome.py` had **70** read calls whereas `fast-pali.cpp` had **11** read calls, then looking at the `time` results we can see that the `real` time for the python program was 0m0.249s and the C++ program was 0m0.111s, C++ program was indeed faster than the python version. Another factor that contributes to the C++ program running faster than the python program is due to the fact that C++ is a compiled language while Python is a interpreted language.